

Explanation of LeNet CNN Forward and Backward Pass

LeNet Architecture Overview

Layer-by-Layer Explanation

1. Convolution Layer

2. Batch Normalization Layer

Explanation of Batch Normalization Layer

Forward Pass Formulas

1. Compute Batch Mean
2. Compute Batch Variance
3. Normalize Inputs
4. Scale and Shift
5. Running Mean and Variance (for Inference)

Backward Pass Formulas

1. Gradient with respect to Beta (β)
2. Gradient with respect to Gamma (γ)
3. Gradient with respect to Normalized Inputs (\hat{x})
4. Gradient with respect to Variance (σ_B^2)
5. Gradient with respect to Mean (μ_B)
6. Gradient with respect to Inputs (x_i)

Implementation Notes

3. ReLU Activation

4. Max Pooling Layer

5. Fully Connected Layer

Complete Forward and Backward Pass for LeNet

Forward Pass:

Backward Pass:

Dropout Layer

Intuitive Idea

Forward Pass

Backward Pass

Summary

The UML class diagram of the LeNet-5 implementation:

Key Points

LeNet Implementation from Scratch

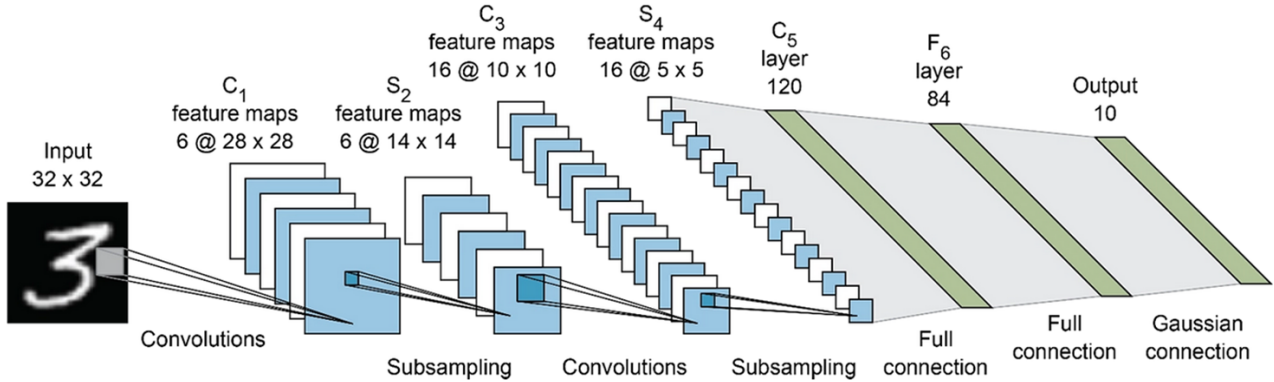
Key Features of This Implementation:

Explanation of LeNet CNN Forward and Backward Pass

LeNet Architecture Overview

LeNet is a classic convolution neural network architecture designed for handwritten digit recognition. It consists of:

1. Convolution layers for feature extraction
2. Activation functions (ReLU) for non-linearity
3. Pooling layers for dimensionality reduction
4. Fully connected layers for classification



Layer-by-Layer Explanation

1. Convolution Layer

Intuitive Idea:

The convolution layer is designed to automatically detect local patterns in the input data, such as edges, textures, or more complex shapes in images. By applying multiple filters, the network can learn hierarchical representations of the input.

Forward Pass:

The convolution operation involves sliding filters over the input to produce feature maps:

$$O_{i,j}^{(c)} = \sum_{m=0}^{f-1} \sum_{n=0}^{f-1} I_{i+m,j+n}^{(k)} \cdot K_{m,n}^{(c,k)} + b^{(c)}$$

where:

- I is the input feature map
- K is the convolution kernel
- b is the bias term
- f is the filter size
- c is the output channel index
- k is the input channel index

Backward Pass:

The gradients are computed using the chain rule:

$$\begin{aligned}\frac{\partial L}{\partial K^{(c,k)}} &= I^{(k)} * \text{rot180} \left(\frac{\partial L}{\partial O^{(c)}} \right) \\ \frac{\partial L}{\partial I^{(k)}} &= \sum_c \text{rot180}(K^{(c,k)}) * \frac{\partial L}{\partial O^{(c)}} \\ \frac{\partial L}{\partial b^{(c)}} &= \sum_{i,j} \frac{\partial L}{\partial O_{i,j}^{(c)}}\end{aligned}$$

2. Batch Normalization Layer

Intuitive Idea:

Batch normalization aims to stabilize and accelerate the training process by normalizing the inputs to each layer, reducing internal covariate shift. It also provides a regularizing effect.

Forward Pass:

Batch normalization normalizes the input, scales, and shifts it:

$$\begin{aligned}\mu_B &= \frac{1}{m} \sum_{i=1}^m x_i \\ \sigma_B^2 &= \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \\ \hat{x}_i &= \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \\ y_i &= \gamma \hat{x}_i + \beta\end{aligned}$$

Backward Pass:

The gradients are computed as:

$$\begin{aligned}\frac{\partial L}{\partial \beta} &= \sum_{i=1}^m \frac{\partial L}{\partial y_i} \\ \frac{\partial L}{\partial \gamma} &= \sum_{i=1}^m \frac{\partial L}{\partial y_i} \hat{x}_i \\ \frac{\partial L}{\partial \hat{x}_i} &= \gamma \frac{\partial L}{\partial y_i} \\ \frac{\partial L}{\partial x_i} &= \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} \left(1 - \frac{1}{m} - \frac{(x_i - \mu_B)^2}{m(\sigma_B^2 + \epsilon)} \right)\end{aligned}$$

Explanation of Batch Normalization Layer

Batch normalization is a technique designed to improve the training of neural networks by normalizing the inputs to each layer. This normalization helps to:

- Reduce internal covariate shift (changes in the distribution of layer inputs during training)
- Accelerate training by allowing higher learning rates
- Provide some regularization effect
- Make the network more robust to initialization

Forward Pass Formulas

1. Compute Batch Mean

The mean of the inputs for each feature map across the batch is calculated:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i$$

where:

- m is the number of samples in the batch
- x_i represents the input for each sample in the batch

2. Compute Batch Variance

The variance of the inputs for each feature map across the batch is calculated:

$$\sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

This measures how much the inputs vary from the mean.

3. Normalize Inputs

The inputs are normalized using the computed mean and variance:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

where:

- ϵ is a small constant (e.g., 10^{-5}) added for numerical stability

4. Scale and Shift

The normalized inputs are scaled and shifted using learnable parameters:

$$y_i = \gamma \hat{x}_i + \beta$$

where:

- γ (scale parameter) and β (shift parameter) are learned during training
- These parameters allow the network to recover the original representation if needed

5. Running Mean and Variance (for Inference)

During training, we maintain running estimates of the mean and variance to use during inference:

$$\mu_{\text{running}} = \text{momentum} \cdot \mu_{\text{running}} + (1 - \text{momentum}) \cdot \mu_B$$

$$\sigma_{\text{running}}^2 = \text{momentum} \cdot \sigma_{\text{running}}^2 + (1 - \text{momentum}) \cdot \sigma_B^2$$

where:

- momentum is typically a value close to 1 (e.g., 0.9)

Backward Pass Formulas

1. Gradient with respect to Beta (β)

The gradient for the shift parameter β is simply the sum of the gradients from the output:

$$\frac{\partial L}{\partial \beta} = \sum_{i=1}^m \frac{\partial L}{\partial y_i}$$

2. Gradient with respect to Gamma (γ)

The gradient for the scale parameter γ involves the normalized inputs:

$$\frac{\partial L}{\partial \gamma} = \sum_{i=1}^m \frac{\partial L}{\partial y_i} \hat{x}_i$$

3. Gradient with respect to Normalized Inputs (\hat{x})

The gradient propagates through the scaling operation:

$$\frac{\partial L}{\partial \hat{x}_i} = \gamma \frac{\partial L}{\partial y_i}$$

4. Gradient with respect to Variance (σ_B^2)

The gradient with respect to the variance involves the deviation of each input from the mean:

$$\frac{\partial L}{\partial \sigma_B^2} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{-1}{2} (x_i - \mu_B) (\sigma_B^2 + \epsilon)^{-3/2}$$

5. Gradient with respect to Mean (μ_B)

The gradient with respect to the mean has two components:

$$\frac{\partial L}{\partial \mu_B} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{-1}{\sqrt{\sigma_B^2 + \epsilon}} + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{-2}{m} \sum_{j=1}^m (x_j - \mu_B)$$

6. Gradient with respect to Inputs (x_i)

The final gradient with respect to the inputs combines the gradients from all previous steps:

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{1}{\sqrt{\sigma_B^2 + \epsilon}} + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{2}{m} (x_i - \mu_B) + \frac{\partial L}{\partial \mu_B} \cdot \frac{1}{m}$$

Implementation Notes

- During training, we use the batch mean and variance for normalization.
- During inference (testing), we use the running mean and variance estimates.
- The parameters γ and β are updated during back propagation like other network weights.
- Batch normalization is typically applied after convolutional or fully connected layers but before activation functions.

3. ReLU Activation

Intuitive Idea:

The Rectified Linear Unit (ReLU) activation function introduces non-linearity to the network while being computationally efficient. It helps the network learn complex patterns by allowing it to model non-linear relationships.

Forward Pass:

The ReLU activation function is applied element-wise:

$$f(x) = \max(0, x)$$

Backward Pass:

The gradient is computed as:

$$\frac{\partial L}{\partial x} = \begin{cases} \frac{\partial L}{\partial y} & \text{if } x > 0 \\ 0 & \text{otherwise} \end{cases}$$

4. Max Pooling Layer

Intuitive Idea:

Max pooling reduces the spatial dimensions of the feature maps, making the network more computationally efficient and helping to prevent overfitting. It also provides a form of translation invariance.

Forward Pass:

Max pooling takes the maximum value within each pooling window:

$$O_{i,j} = \max_{m,n \in \text{pooling window}} I_{i+m,j+n}$$

Backward Pass:

The gradient is propagated only to the positions of the maximum values:

$$\frac{\partial L}{\partial I_{i,j}} = \begin{cases} \frac{\partial L}{\partial O_{k,l}} & \text{if } I_{i,j} = \max \text{ in window} \\ 0 & \text{otherwise} \end{cases}$$

5. Fully Connected Layer

Intuitive Idea:

Fully connected layers take the high-level features extracted by the convolutional layers and combine them to make predictions. They perform classification based on the learned features.

Forward Pass:

The fully connected layer performs matrix multiplication:

$$y = Wx + b$$

Backward Pass:

The gradients are computed using matrix operations:

$$\frac{\partial L}{\partial W} = x \cdot \frac{\partial L}{\partial y}^T$$

$$\frac{\partial L}{\partial b} = \sum_{i=1}^m \frac{\partial L}{\partial y_i}$$
$$\frac{\partial L}{\partial x} = W^T \cdot \frac{\partial L}{\partial y}$$

Complete Forward and Backward Pass for LeNet

Forward Pass:

1. Input image passes through the first convolutional layer
2. Batch normalization is applied to the output
3. ReLU activation introduces non-linearity
4. Max pooling reduces spatial dimensions
5. The process repeats for subsequent layers
6. Flattened output is passed through fully connected layers
7. Softmax activation produces probability distribution for classification

Backward Pass:

1. Gradients start from the output layer (usually computed using cross-entropy loss)
 2. Gradients propagate backward through fully connected layers
 3. Max pooling layer propagates gradients only to max positions
 4. ReLU layer passes gradients where activations were positive
 5. Batch normalization layer updates its parameters (gamma and beta)
 6. Convolution layers update their filters and biases
 7. Gradients continue propagating backward through all layers
-

Dropout Layer

Intuitive Idea

The dropout layer is a regularization technique used in neural networks to prevent overfitting. The core idea is to randomly "drop out" a certain percentage of neurons during the training phase. This forces the network to learn more robust features and prevents it from relying too much on any single neuron, thereby improving generalization.

Forward Pass

During the forward pass, each neuron has a probability p of being dropped. The remaining neurons' outputs are scaled by a factor of $1/(1-p)$ to maintain the expected sum of outputs. This scaling ensures that the magnitudes of the activations remain consistent across training and inference phases.

Formula:

For a neuron's output a , the forward pass with dropout can be represented as:

```
1 | a' = (a * mask) / (1 - p)
```

where mask is a binary tensor of the same shape as \mathbf{a} , with elements set to 1 with probability $(1-p)$ and 0 with probability p .

Backward Pass

During the backward pass, gradients are only propagated through the neurons that were not dropped out during the forward pass. The gradients for the dropped neurons are zeroed out. This means that only the active neurons (those not dropped) contribute to the weight updates.

Formula:

If δ_{delta} represents the gradient of the loss with respect to the output of the dropout layer, then the gradient with respect to the input of the dropout layer is:

```
1 delta_input = delta * mask / (1 - p)
```


This ensures that the gradients are scaled appropriately and only propagate through the active neurons.

Summary

- **Intuitive Idea:** Randomly deactivate neurons to prevent over-reliance on any single neuron.
- **Forward Pass:** Apply dropout mask and scale outputs.
- **Backward Pass:** Propagate gradients only through active neurons and scale appropriately.

Dropout is particularly effective in fully connected layers and is often used in conjunction with other regularization techniques like L2 regularization.

The UML class diagram of the LeNet-5 implementation:




```

5 | + forward(x: np.ndarray) : np.ndarray |
6 | + backward(grad: np.ndarray) : np.ndarray |
7 | + update(learning_rate: float) : void |
8 +-----+
9 | ^ |
10 | | |
11 | | |
12 +-----+
13 | | | | |
14 | | | | |
15 | +-----+ +-----+ +-----+
16 | | Conv2D | | MaxPool2D | | Dense |
17 | +-----+ +-----+ +-----+
18 | | - in_channels: int | | - pool_size: int | | - input_dim: int |
19 | | - out_channels: int | | - stride: int | | - output_dim: int |
20 | | - kernel_size: int | +-----+ +-----+
21 | | - stride: int | | + forward() | | + forward() |
22 | | - padding: int | | + backward() | | + backward() |
23 | | - weights: np.ndarray | +-----+ +-----+
24 | | - biases: np.ndarray |
25 | | - cache: tuple |
26 | +-----+
27 | | + im2col() |
28 | | + col2im() |
29 | | + forward() |
30 | | + backward() |
31 | | + update() |
32 | +-----+
33 |
34 | +-----+ +-----+ +-----+
35 | | BatchNorm2D | | Dropout | | ReLU |
36 | +-----+ +-----+ +-----+
37 | | - gamma: np.ndarray | | - p: float | | - cache: np.ndarray |
38 | | - beta: np.ndarray | | - mask: np.ndarray | +-----+
39 | | - eps: float | +-----+ | + forward() |
40 | | - momentum: float | | + forward() | | + backward() |
41 | | - running_mean: np.ndarray | + backward() | +-----+
42 | | - running_var: np.ndarray |
43 | | - cache: tuple | +-----+ +-----+
44 | +-----+ | Softmax |
45 | | + forward() | +-----+
46 | | + backward() | | - cache: np.ndarray |
47 | | + update() | +-----+
48 | +-----+ | + forward() |
49 | | + backward() |
50 | +-----+ +-----+
51 | | Flatten |
52 | +-----+
53 | | - input_shape: tuple |
54 | +-----+
55 | | + forward() |
56 | | + backward() |
57 | +-----+
58 |
59 +-----+
60 | LeNet5 |
61 +-----+
62 | - layers: List[Layer] |

```

```

63 | +-----+
64 | + forward(x: np.ndarray, training: bool) : np.ndarray |
65 | + backward(grad: np.ndarray) : void |
66 | + update(learning_rate: float) : void |
67 | + predict(x: np.ndarray) : np.ndarray |
68 | + loss(y_pred, y_true) : float |
69 | + accuracy(y_pred, y_true) : float |
70 | +-----+

```

Key Points

Each concrete layer implements:

- `forward()`: Computes the forward pass
- `backward()`: Computes gradients during back propagation
- `update()`: Updates parameters (for layers with learnable weights)
- All layers maintain their own parameters and cache variables needed for back propagation.

The `LeNet5` class orchestrates:

- Full forward/backward passes through all layers
- Parameter updates
- Prediction and evaluation methods

LeNet Implementation from Scratch

A complete implementation of LeNet-5 using NumPy, featuring:

- Convolution layers with `im2col` optimization
- Max pooling layers
- Fully connected layers
- Batch normalization
- Dropout
- ReLU activation
- Softmax output

```

1 | import numpy as np
2 | from typing import Tuple, List, Optional
3 |
4 | class Layer:
5 |     def forward(self, x: np.ndarray) -> np.ndarray:

```

```

6         raise NotImplementedError
7
8     def backward(self, grad: np.ndarray) -> np.ndarray:
9         raise NotImplementedError
10
11     def update(self, learning_rate: float) -> None:
12         pass
13
14 import numpy as np
15 from numpy.lib.stride_tricks import as_strided
16 from typing import Tuple, List, Optional
17
18 class Layer:
19     def forward(self, x: np.ndarray) -> np.ndarray:
20         raise NotImplementedError
21
22     def backward(self, grad: np.ndarray) -> np.ndarray:
23         raise NotImplementedError
24
25     def update(self, learning_rate: float) -> None:
26         pass
27
28 class Conv2D(Layer):
29     def __init__(self, in_channels: int, out_channels: int, kernel_size: int,
30                 stride: int = 1, padding: int = 0):
31         self.in_channels = in_channels
32         self.out_channels = out_channels
33         self.kernel_size = kernel_size
34         self.stride = stride
35         self.padding = padding
36
37         # He initialization
38         scale = np.sqrt(2.0 / (in_channels * kernel_size * kernel_size))
39         self.weights = np.random.randn(out_channels, in_channels, kernel_size, kernel_size)
40         * scale
41         self.biases = np.zeros(out_channels)
42
43         self.cache = None
44
45     def im2col(self, x: np.ndarray) -> np.ndarray:
46         """Optimized im2col using as_strided"""
47         N, C, H, W = x.shape
48         K = self.kernel_size
49         stride = self.stride
50         pad = self.padding
51
52         # Add padding
53         if pad > 0:
54             x_padded = np.pad(x, ((0, 0), (0, 0), (pad, pad), (pad, pad)), mode='constant')
55         else:
56             x_padded = x
57
58         # Calculate output dimensions
59         H_out = (H + 2 * pad - K) // stride + 1
60         W_out = (W + 2 * pad - K) // stride + 1
61
62         # Shape of the output array
63         shape = (N, C, K, K, H_out, W_out)

```

```

63
64     # Strides of the input array (bytes to move in each dimension)
65     strides = (x_padded.strides[0], x_padded.strides[1],
66               x_padded.strides[2], x_padded.strides[3],
67               stride * x_padded.strides[2], stride * x_padded.strides[3])
68
69     # Create strided view
70     strided = as_strided(x_padded, shape=shape, strides=strides)
71
72     # Reshape to column format
73     cols = strided.transpose(0, 4, 5, 1, 2, 3).reshape(N * H_out * W_out, -1)
74
75     return cols.T # Transpose to match standard im2col output
76
77 def col2im(self, cols: np.ndarray, x_shape: Tuple[int, int, int, int]) -> np.ndarray:
78     """Inverse of im2col using numpy operations"""
79     N, C, H, W = x_shape
80     K = self.kernel_size
81     stride = self.stride
82     pad = self.padding
83
84     H_padded = H + 2 * pad
85     W_padded = W + 2 * pad
86     x_padded = np.zeros((N, C, H_padded, W_padded), dtype=cols.dtype)
87
88     H_out = (H + 2 * pad - K) // stride + 1
89     W_out = (W + 2 * pad - K) // stride + 1
90
91     # Reshape columns back to image patches
92     cols_resaped = cols.T.reshape(N, H_out, W_out, C, K, K)
93     cols_resaped = cols_resaped.transpose(0, 3, 4, 5, 1, 2) # N, C, K, K, H_out,
W_out
94
95     # Add patches back to image
96     for i in range(K):
97         for j in range(K):
98             x_padded[:, :, i:H_out*stride:stride, j:W_out*stride:stride] +=
cols_resaped[:, :, i, j, :, :]
99
100     if pad == 0:
101         return x_padded
102     return x_padded[:, :, pad:-pad, pad:-pad]
103
104 def forward(self, x: np.ndarray) -> np.ndarray:
105     N, C, H, W = x.shape
106     K = self.kernel_size
107     stride = self.stride
108     pad = self.padding
109
110     H_out = (H + 2 * pad - K) // stride + 1
111     W_out = (W + 2 * pad - K) // stride + 1
112
113     # Convert input to columns using optimized im2col
114     x_cols = self.im2col(x)
115
116     # Reshape weights
117     w_cols = self.weights.reshape(self.out_channels, -1)
118

```

```

119         # Perform matrix multiplication
120         out = w_cols @ x_cols + self.biases.reshape(-1, 1)
121
122         # Reshape output
123         out = out.reshape(self.out_channels, H_out, W_out, N)
124         out = out.transpose(3, 0, 1, 2)
125
126         # Cache for backward pass
127         self.cache = (x, x_cols)
128
129         return out
130
131     def backward(self, grad: np.ndarray) -> np.ndarray:
132         x, x_cols = self.cache
133         N, C, H, W = x.shape
134         K = self.kernel_size
135
136         # Reshape gradient (outc, oh, hw, n) => (outc, oh*ow*n)
137         grad_resaped = grad.transpose(1, 2, 3, 0).reshape(self.out_channels, -1)
138
139         # Calculate weight gradients
140         #(outc, oh*ow*n) @ (oh*ow*n, in_c*k*k)
141         dw = grad_resaped @ x_cols.T
142         # (outc, in_c, k, k)
143         dw = dw.reshape(self.weights.shape)
144
145         # Calculate bias gradients db(outc, 1)
146         db = np.sum(grad, axis=(0, 2, 3))
147
148         # Calculate input gradients, w_cols(out_c, in_c*k*k)
149         w_cols = self.weights.reshape(self.out_channels, -1)
150         dx_cols = w_cols.T @ grad_resaped
151         dx = self.col2im(dx_cols, x.shape)
152
153         # Store gradients
154         self.dw = dw
155         self.db = db
156
157         return dx
158
159     def update(self, learning_rate: float) -> None:
160         self.weights -= learning_rate * self.dw
161         self.biases -= learning_rate * self.db
162
163     #
164     # The other classes (MaxPool2D, Flatten, Dense, BatchNorm2D, Dropout, ReLU, Softmax,
165     # LeNet5)
166     #
167     class MaxPool2D(Layer):
168         def __init__(self, pool_size: int = 2, stride: int = 2):
169             self.pool_size = pool_size
170             self.stride = stride
171             self.cache = None
172
173         def forward(self, x: np.ndarray) -> np.ndarray:
174             N, C, H, W = x.shape
175             pool_size = self.pool_size
176             stride = self.stride

```

```

176
177     H_out = (H - pool_size) // stride + 1
178     W_out = (W - pool_size) // stride + 1
179
180     # Reshape input for vectorized max operation
181     x_resaped = x.reshape(N, C, H // pool_size, pool_size,
182                           W // pool_size, pool_size)
183     out = x_resaped.max(axis=3).max(axis=4)
184
185     # Create mask for backward pass
186     x_resaped = x.reshape(N, C, H_out, stride, W_out, stride)
187     mask = (x_resaped == out[:, :, :, np.newaxis, :, np.newaxis])
188
189     self.cache = (x.shape, mask)
190     return out
191
192     def backward(self, grad: np.ndarray) -> np.ndarray:
193         input_shape, mask = self.cache
194         N, C, H, W = input_shape
195         pool_size = self.pool_size
196
197         # Upsample gradient
198         grad_upsampled = np.repeat(np.repeat(grad, pool_size, axis=2), pool_size, axis=3)
199
200         # Apply mask
201         dx = grad_upsampled * mask.reshape(input_shape)
202         return dx
203
204     class Flatten(Layer):
205         def __init__(self):
206             self.input_shape = None
207
208         def forward(self, x: np.ndarray) -> np.ndarray:
209             self.input_shape = x.shape
210             return x.reshape(x.shape[0], -1)
211
212         def backward(self, grad: np.ndarray) -> np.ndarray:
213             return grad.reshape(self.input_shape)
214
215     class Dense(Layer):
216         def __init__(self, input_dim: int, output_dim: int):
217             # He initialization
218             scale = np.sqrt(2.0 / input_dim)
219             self.weights = np.random.randn(input_dim, output_dim) * scale
220             self.biases = np.zeros(output_dim)
221             self.cache = None
222
223         def forward(self, x: np.ndarray) -> np.ndarray:
224             self.cache = x
225             return x @ self.weights + self.biases
226
227         def backward(self, grad: np.ndarray) -> np.ndarray:
228             x = self.cache
229             self.dw = x.T @ grad
230             self.db = np.sum(grad, axis=0)
231             return grad @ self.weights.T
232
233         def update(self, learning_rate: float) -> None:

```

```

234         self.weights -= learning_rate * self.dw
235         self.biases -= learning_rate * self.db
236
237     class BatchNorm2D(Layer):
238         def __init__(self, num_features: int, eps: float = 1e-5, momentum: float = 0.1):
239             self.gamma = np.ones(num_features)
240             self.beta = np.zeros(num_features)
241             self.eps = eps
242             self.momentum = momentum
243             self.running_mean = np.zeros(num_features)
244             self.running_var = np.ones(num_features)
245             self.cache = None
246
247         def forward(self, x: np.ndarray, training: bool = True) -> np.ndarray:
248             N, C, H, W = x.shape
249
250             if training:
251                 # Calculate mean and variance over batch and spatial dimensions
252                 mean = np.mean(x, axis=(0, 2, 3), keepdims=True)
253                 var = np.var(x, axis=(0, 2, 3), keepdims=True)
254
255                 # Update running statistics
256                 self.running_mean = self.momentum * mean.squeeze() + (1 - self.momentum) *
self.running_mean
257                 self.running_var = self.momentum * var.squeeze() + (1 - self.momentum) *
self.running_var
258
259                 # Normalize
260                 x_norm = (x - mean) / np.sqrt(var + self.eps)
261             else:
262                 # Use running statistics during inference
263                 x_norm = (x - self.running_mean.reshape(1, C, 1, 1)) /
np.sqrt(self.running_var.reshape(1, C, 1, 1) + self.eps)
264
265             # Scale and shift
266             out = self.gamma.reshape(1, C, 1, 1) * x_norm + self.beta.reshape(1, C, 1, 1)
267
268             if training:
269                 self.cache = (x, mean, var, x_norm)
270
271             return out
272
273         def backward(self, grad: np.ndarray) -> np.ndarray:
274             x, mean, var, x_norm = self.cache
275             N, C, H, W = x.shape
276
277             # Calculate gradients
278             dbeta = np.sum(grad, axis=(0, 2, 3))
279             dgamma = np.sum(grad * x_norm, axis=(0, 2, 3))
280
281             # Intermediate gradients
282             dx_norm = grad * self.gamma.reshape(1, C, 1, 1)
283             dvar = np.sum(dx_norm * (x - mean) * -0.5 * (var + self.eps) ** (-1.5), axis=(0, 2,
3), keepdims=True)
284             dmean = np.sum(dx_norm * -1 / np.sqrt(var + self.eps), axis=(0, 2, 3),
keepdims=True) + \
285                 dvar * np.mean(-2 * (x - mean), axis=(0, 2, 3), keepdims=True))
286

```

```

287     # Final gradient
288     dx = dx_norm / np.sqrt(var + self.eps) + \
289         dvar * 2 * (x - mean) / (N * H * W) + \
290         dmean / (N * H * W)
291
292     self.dgamma = dgamma
293     self.dbeta = dbeta
294
295     return dx
296
297     def update(self, learning_rate: float) -> None:
298         self.gamma -= learning_rate * self.dgamma
299         self.beta -= learning_rate * self.dbeta
300
301     # inverted drop-out implementation
302     class Dropout(Layer):
303         def __init__(self, p: float = 0.5):
304             self.p = p
305             self.mask = None
306
307         def forward(self, x: np.ndarray, training: bool = True) -> np.ndarray:
308             if training:
309                 self.mask = (np.random.rand(*x.shape) > self.p) / (1 - self.p)
310                 return x * self.mask
311             return x
312
313         def backward(self, grad: np.ndarray) -> np.ndarray:
314             return grad * self.mask
315
316     class ReLU(Layer):
317         def __init__(self):
318             self.cache = None
319
320         def forward(self, x: np.ndarray) -> np.ndarray:
321             self.cache = x
322             return np.maximum(0, x)
323
324         def backward(self, grad: np.ndarray) -> np.ndarray:
325             x = self.cache
326             return grad * (x > 0)
327
328     class Softmax(Layer):
329         def __init__(self):
330             self.cache = None
331
332         def forward(self, x: np.ndarray) -> np.ndarray:
333             # Numerically stable softmax
334             exps = np.exp(x - np.max(x, axis=1, keepdims=True))
335             out = exps / np.sum(exps, axis=1, keepdims=True)
336             self.cache = out
337             return out
338
339         def backward(self, grad: np.ndarray) -> np.ndarray:
340             s = self.cache
341             # Jacobian matrix: diag(s) - s.T @ s
342             return s * (grad - (grad * s).sum(axis=1, keepdims=True))
343
344     class LeNet5:

```



```

345 def __init__(self, input_shape: Tuple[int, int, int], num_classes: int):
346     C, H, W = input_shape
347
348     self.layers = [
349         Conv2D(in_channels=C, out_channels=6, kernel_size=5, stride=1, padding=2),
350         BatchNorm2D(num_features=6),
351         ReLU(),
352         MaxPool2D(pool_size=2, stride=2),
353
354         Conv2D(in_channels=6, out_channels=16, kernel_size=5, stride=1, padding=0),
355         BatchNorm2D(num_features=16),
356         ReLU(),
357         MaxPool2D(pool_size=2, stride=2),
358
359         Flatten(),
360
361         Dense(input_dim=16*5*5, output_dim=120),
362         BatchNorm2D(num_features=120),
363         ReLU(),
364         Dropout(p=0.5),
365
366         Dense(input_dim=120, output_dim=84),
367         BatchNorm2D(num_features=84),
368         ReLU(),
369         Dropout(p=0.5),
370
371         Dense(input_dim=84, output_dim=num_classes),
372         Softmax()
373     ]
374
375 def forward(self, x: np.ndarray, training: bool = True) -> np.ndarray:
376     for layer in self.layers:
377         if isinstance(layer, (BatchNorm2D, Dropout)):
378             x = layer.forward(x, training=training)
379         else:
380             x = layer.forward(x)
381     return x
382
383 def backward(self, grad: np.ndarray) -> None:
384     for layer in reversed(self.layers):
385         grad = layer.backward(grad)
386
387 def update(self, learning_rate: float) -> None:
388     for layer in self.layers:
389         if hasattr(layer, 'update'):
390             layer.update(learning_rate)
391
392 def predict(self, x: np.ndarray) -> np.ndarray:
393     return np.argmax(self.forward(x, training=False), axis=1)
394
395 def loss(self, y_pred: np.ndarray, y_true: np.ndarray) -> float:
396     # Cross-entropy loss
397     m = y_true.shape[0]
398     log_likelihood = -np.log(y_pred[range(m), y_true])
399     loss = np.sum(log_likelihood) / m
400     return loss
401
402 def accuracy(self, y_pred: np.ndarray, y_true: np.ndarray) -> float:

```

```

403         return np.mean(y_pred == y_true)
404
405     # Example usage:
406     if __name__ == "__main__":
407         # Create a dummy dataset
408         np.random.seed(42)
409         X_train = np.random.randn(100, 1, 28, 28) # 100 samples, 1 channel, 28x28
410         y_train = np.random.randint(0, 10, size=100) # 10 classes
411
412         # Initialize LeNet-5
413         lenet = LeNet5(input_shape=(1, 28, 28), num_classes=10)
414
415         # Training loop (simplified)
416         learning_rate = 0.01
417         epochs = 5
418         batch_size = 10
419
420         for epoch in range(epochs):
421             epoch_loss = 0
422             correct = 0
423
424             for i in range(0, len(X_train), batch_size):
425                 X_batch = X_train[i:i+batch_size]
426                 y_batch = y_train[i:i+batch_size]
427
428                 # Forward pass
429                 y_pred = lenet.forward(X_batch)
430
431                 # Calculate loss
432                 loss = lenet.loss(y_pred, y_batch)
433                 epoch_loss += loss * len(X_batch)
434
435                 # Calculate accuracy
436                 preds = np.argmax(y_pred, axis=1)
437                 correct += np.sum(preds == y_batch)
438
439                 # Backward pass
440                 grad = y_pred.copy()
441                 grad[len(y_batch), y_batch] -= 1
442                 grad /= len(y_batch)
443                 lenet.backward(grad)
444
445                 # Update weights
446                 lenet.update(learning_rate)
447
448             epoch_loss /= len(X_train)
449             accuracy = correct / len(X_train)
450
451             print(f"Epoch {epoch+1}/{epochs}, Loss: {epoch_loss:.4f}, Accuracy:
{accuracy:.4f}")

```

Key Features of This Implementation:

1. Convolutional Layers with im2col:

- Uses the im2col algorithm for efficient convolution operations

- Implements both forward and backward passes
- Supports padding and stride

2. **Batch Normalization:**

- Implements batch norm for both convolutional and fully connected layers
- Maintains running statistics for inference
- Properly handles the scale (gamma) and shift (beta) parameters

3. **Dropout:**

- Implements inverted dropout
- Only active during training
- Scales activations during training to maintain expected values

4. **Complete LeNet-5 Architecture:**

- Two convolutional layers with ReLU and max pooling
- Three fully connected layers
- Batch normalization and dropout for regularization
- Softmax output for classification

5. **Training Framework:**

- Includes forward and backward passes
- Weight updates with learning rate
- Cross-entropy loss calculation
- Accuracy measurement

This implementation provides a complete, from-scratch version of LeNet-5 with modern improvements like batch normalization and dropout, while maintaining the original architecture's essence.
